

A High-Throughput Hardware NoC for Llama Acceleration Tensors

Paper ID Reference: PAP-2026-007

Authors: Contact Author 7, Co-Author 8

Abstract: This research designs a custom Network-on-Chip architecture specifically optimized for sparse attention weight tensors in transformer models.

1. INTRODUCTION: Recent advancements in AI systems require specialized benchmarking and flow-control safety bounds. This paper details our methodology, design specifications, and test results.
2. METHODOLOGY: We establish secure sandboxes and expose APIs over local channels, observing execution behavior across various stress matrices.
3. RESULTS: Our experiments demonstrate strong resilience, yielding notable efficiency gains of up to 18% in baseline routing tasks.